

# hw6-2-muse

February 13, 2025

## 1 HOMEWORK 6: TEXT CLASSIFICATION

In this homework, you will create models to classify texts from TRUE call-center. There are two classification tasks: 1. Action Classification: Identify which action the customer would like to take (e.g. enquire, report, cancel) 2. Object Classification: Identify which object the customer is referring to (e.g. payment, true money, internet, roaming)

We will focus only on the Object Classification task for this homework.

In this homework, you are asked to compare different text classification models in terms of accuracy and inference time.

You will need to build 3 different models.

1. A model based on tf-idf
2. A model based on MUSE
3. A model based on wangchanBERTa

**You will be asked to submit 3 different files (.pdf from .ipynb) that do the 3 different models. Finally, answer the accuracy and runtime numbers in MCV.**

This homework is quite free form, and your answer may vary. We hope that the processing during the course of this assignment will make you think more about the design choices in text classification.

```
[1]: # !wget --no-check-certificate https://www.dropbox.com/s/37u83g55p19kvrl/  
      ↪ clean-phone-data-for-students.csv
```

```
[2]: %pip install pythainlp
```

```
Requirement already satisfied: pythainlp in  
/Users/pupipatsingkhorn/miniconda3/envs/datascience/lib/python3.11/site-packages  
(5.0.5)  
Requirement already satisfied: requests>=2.22.0 in  
/Users/pupipatsingkhorn/miniconda3/envs/datascience/lib/python3.11/site-packages  
(from pythainlp) (2.32.3)  
Requirement already satisfied: charset-normalizer<4,>=2 in  
/Users/pupipatsingkhorn/miniconda3/envs/datascience/lib/python3.11/site-packages  
(from requests>=2.22.0->pythainlp) (3.3.2)  
Requirement already satisfied: idna<4,>=2.5 in  
/Users/pupipatsingkhorn/miniconda3/envs/datascience/lib/python3.11/site-packages  
(from requests>=2.22.0->pythainlp) (3.10)  
Requirement already satisfied: urllib3<3,>=1.21.1 in
```

```

/Users/pupipatsingkhorn/miniconda3/envs/datascience/lib/python3.11/site-packages
(from requests>=2.22.0->pythainlp) (2.3.0)
Requirement already satisfied: certifi>=2017.4.17 in
/Users/pupipatsingkhorn/miniconda3/envs/datascience/lib/python3.11/site-packages
(from requests>=2.22.0->pythainlp) (2025.1.31)
Note: you may need to restart the kernel to use updated packages.

```

## 1.1 Import Libs

```

[3]: %matplotlib inline
import pandas
import sklearn
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd

from torch.utils.data import Dataset
from IPython.display import display
from collections import defaultdict
from sklearn.metrics import accuracy_score

import warnings
warnings.filterwarnings('ignore')
import time

```

```

[4]: SEED = 42

```

## 1.2 Loading data

First, we load the data from disk into a Dataframe.

A Dataframe is essentially a table, or 2D-array/Matrix with a name for each column.

```

[5]: data_df = pd.read_csv('clean-phone-data-for-students.csv')

```

Let's preview the data.

```

[6]: # Show the top 5 rows
display(data_df.head())
# Summarize the data
data_df.describe()

```

			Sentence	Utterance	Action	Object
0	<PHONE_NUMBER_REMOVED>		Counte...	enquire	payment	
1	internet			enquire	package	
2	...	report		suspend		
3	internet	...	enquire	internet		
4	...	report	phone_issues			

```
[6]:
```

	Sentence	Utterance	Action	Object
count		16175	16175	16175
unique		13389	10	33
top		enquire	service	
freq		97	10377	2525

### 1.3 Data cleaning

We call the `DataFrame.describe()` again. Notice that there are 33 unique labels/classes for object and 10 unique labels for action that the model will try to predict. But there are unwanted duplications e.g. `Idd,idd,lotalty_card,Lotalty_card`

Also note that, there are 13389 unique sentence utterances from 16175 utterances. You have to clean that too!

### 1.4 #TODO 0.1:

- You will have to remove unwanted label duplications as well as duplications in text inputs.
- Also, you will have to trim out unwanted whitespaces from the text inputs.

This shouldn't be too hard, as you have already seen it in the demo.

```
[7]: display(data_df.describe())
display(data_df.Object.unique())
display(data_df.Action.unique())
```

	Sentence	Utterance	Action	Object
count		16175	16175	16175
unique		13389	10	33
top		enquire	service	
freq		97	10377	2525

```
array(['payment', 'package', 'suspend', 'internet', 'phone_issues',
      'service', 'nonTrueMove', 'balance', 'detail', 'bill', 'credit',
      'promotion', 'mobile_setting', 'iservice', 'roaming', 'truemoney',
      'information', 'lost_stolen', 'balance_minutes', 'idd',
      'TrueMoney', 'garbage', 'Payment', 'IDD', 'ringtone', 'Idd',
      'rate', 'loyalty_card', 'contact', 'officer', 'Balance', 'Service',
      'Loyalty_card'], dtype=object)

array(['enquire', 'report', 'cancel', 'Enquire', 'buy', 'activate',
      'request', 'Report', 'garbage', 'change'], dtype=object)
```

```
[8]: # TODO 1: Data Cleaning

# Filter cols
cols = ["Sentence Utterance", "Object"]
data_df = data_df[cols]
data_df.columns = ['input', 'raw_label']
```

```

# Lowercase: label
data_df['clean_label']=data_df['raw_label'].str.lower().copy()
data_df.drop('raw_label', axis=1, inplace=True)

# Trim white spaces: input
data_df['input'] = data_df['input'].str.strip()

# Remove duplicate: input
data_df = data_df.drop_duplicates(subset=['input'], keep='first')

# Display summary
display(data_df.describe())
display(data_df['clean_label'].unique())

```

		input	clean_label
count		13367	13367
unique		13367	26
top	<PHONE_NUMBER_REMOVED>	Counter...	service
freq		1	2108

```

array(['payment', 'package', 'suspend', 'internet', 'phone_issues',
      'service', 'nontruemove', 'balance', 'detail', 'bill', 'credit',
      'promotion', 'mobile_setting', 'iservice', 'roaming', 'truemoney',
      'information', 'lost_stolen', 'balance_minutes', 'idd', 'garbage',
      'ringtone', 'rate', 'loyalty_card', 'contact', 'officer'],
      dtype=object)

```

Split data into train, validation, and test sets (normally the ratio will be 80:10:10 , respectively). We recommend to use `train_test_split` from `scikit-learn` to split the data into train, validation, test set.

In addition, it should split the data that distribution of the labels in train, validation, test set are similar. There is **stratify** option to handle this issue.

[https://scikit-learn.org/stable/modules/generated/sklearn.model\\_selection.train\\_test\\_split.html](https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.train_test_split.html)

Make sure the same data splitting is used for all models.

```

[9]: # Mapping
data = data_df.to_numpy()

unique_label = data_df.clean_label.unique()

label_2_num_map = dict(zip(unique_label, range(len(unique_label))))
num_2_label_map = dict(zip(range(len(unique_label)), unique_label))

print("Create Mappings")
display(num_2_label_map)
display(label_2_num_map)

```

```

print("Before Mappings")
display(data[:, 1])
data[:,1] = np.vectorize(label_2_num_map.get)(data[:,1]) # Mapping...
print("After Mappings")
display(data[:, 1])

# Trim
def strip_str(string):
    return string.strip()
print("Before")
print(data)
data[:,0] = np.vectorize(strip_str)(data[:,0]) # Trimming...
print("After")
print(data)

```

Create Mappings

```

{0: 'payment',
 1: 'package',
 2: 'suspend',
 3: 'internet',
 4: 'phone_issues',
 5: 'service',
 6: 'nontruemove',
 7: 'balance',
 8: 'detail',
 9: 'bill',
10: 'credit',
11: 'promotion',
12: 'mobile_setting',
13: 'iservice',
14: 'roaming',
15: 'truemoney',
16: 'information',
17: 'lost_stolen',
18: 'balance_minutes',
19: 'idd',
20: 'garbage',
21: 'ringtone',
22: 'rate',
23: 'loyalty_card',
24: 'contact',
25: 'officer'}

{'payment': 0,
 'package': 1,
 'suspend': 2,
 'internet': 3,
 'phone_issues': 4,

```

```

'service': 5,
'nontruemove': 6,
'balance': 7,
'detail': 8,
'bill': 9,
'credit': 10,
'promotion': 11,
'mobile_setting': 12,
'iservice': 13,
'roaming': 14,
'truemoney': 15,
'information': 16,
'lost_stolen': 17,
'balance_minutes': 18,
'idd': 19,
'garbage': 20,
'ringtone': 21,
'rate': 22,
'loyalty_card': 23,
'contact': 24,
'officer': 25}

```

Before Mappings

```

array(['payment', 'package', 'suspend', ..., 'balance', 'balance',
      'package'], dtype=object)

```

After Mappings

```

array([0, 1, 2, ..., 7, 7, 1], dtype=object)

```

Before

```

[['<PHONE_NUMBER_REMOVED>      Counter Services      3276.25
  3057.79      '
  0]
['internet      ' 1]
['      ' 2]
...
['      ' 7]
['      ' 7]
['      ' 1]]

```

After

```

[['<PHONE_NUMBER_REMOVED>      Counter Services      3276.25
  3057.79      '
  0]
['internet      ' 1]
['      ' 2]
...
['      ' 7]
['      ' 7]

```

```
[' ' 1]]
```

```
[10]: # TODO: Split data
from sklearn.model_selection import train_test_split

SEED = 42

def split_data(data_df, random_state=SEED):
    """split_data splits the data into train:validation:test=8:1:1 sets."""

    def _filter_data(data_df):
        X = data_df["input"]
        y = data_df["clean_label"]
        # Drop classes with fewer than 10(8:1:1) instances
        class_counts = y.value_counts()
        valid_classes = class_counts[class_counts >= 10].index
        filtered_data = data_df[data_df["clean_label"].isin(valid_classes)]
        # Update X and y after filtering
        X = filtered_data["input"]
        y = filtered_data["clean_label"]
        return X, y.astype(int)

    X, y = _filter_data(data_df)

    # First split: Train (80%) and Temp (20%)
    X_train, X_temp, y_train, y_temp = train_test_split(
        X, y, test_size=0.20, stratify=y, random_state=random_state
    )

    # Second split: Validation (10%) and Test (10%)
    X_val, X_test, y_val, y_test = train_test_split(
        X_temp, y_temp, test_size=0.50, stratify=y_temp,
        random_state=random_state
    )

    # Display dataset sizes
    print(f"Train size: {len(X_train)}")
    print(f"Validation size: {len(X_val)}")
    print(f"Test size: {len(X_test)}")

    return X_train, X_val, X_test, y_train, y_val, y_test

# Split
df = pd.DataFrame(data, columns=['input', 'clean_label'])
X_train, X_val, X_test, y_train, y_val, y_test = split_data(df)
```

Train size: 10690

Validation size: 1336

Test size: 1337

## 2 Model 2 MUSE

Build a simple logistic regression model using features from the MUSE model.

Which MUSE model will you use? Why?

**Ans:** `sentence-transformers/use-cmlm-multilingual` because - Pre-trained on multiple languages, including Thai - Captures sentence semantics - Achieves better generalization compared to traditional vector-based models like TF-IDF

MUSE is typically used with tensorflow. However, there are some pytorch conversions made by some people.

<https://huggingface.co/sentence-transformers/use-cmlm-multilingual>

<https://huggingface.co/dayyass/universal-sentence-encoder-multilingual-large-3-pytorch>

```
[11]: # %pip install -U sentence-transformers
      # %pip install tf-keras
```

```
[12]: from sentence_transformers import SentenceTransformer
      from sklearn.linear_model import LogisticRegression
      start_time = time.time()
      print("MUSE + Logistic Regression")

      # Load the pre-trained MUSE model from HuggingFace
      muse_model = SentenceTransformer("sentence-transformers/use-cmlm-multilingual")

      # Function to encode text using MUSE
      def encode_texts(texts):
          return muse_model.encode(texts, convert_to_numpy=True)

      # Encode training, validation, and test sets
      start_enc_time = time.time()
      X_train_enc = encode_texts(X_train.tolist())
      X_val_enc = encode_texts(X_val.tolist())
      X_test_enc = encode_texts(X_test.tolist())
      end_enc_time = time.time()
      print(f"Encoding Time: {end_enc_time - start_enc_time:.4f} seconds")

      # Logistic Regression model
      model = LogisticRegression(random_state=SEED)

      # Train the model
      start_train_time = time.time()
      model.fit(X_train_enc, y_train) # training...
      end_train_time = time.time()
```



```

print(f"Training Time: {end_train_time - start_train_time:.4f} seconds")

# Predictions
y_pred_train = model.predict(X_train_enc)
y_pred_val = model.predict(X_val_enc)
y_pred_test = model.predict(X_test_enc)

# Evaluate model accuracy
train_acc = np.mean(y_train.astype(int) == y_pred_train)
val_acc = np.mean(y_val.astype(int) == y_pred_val)
test_acc = np.mean(y_test.astype(int) == y_pred_test)

print(f"Train Accuracy: {train_acc:.4f}")
print(f"Validation Accuracy: {val_acc:.4f}")
print(f"Test Accuracy: {test_acc:.4f}")
end_time = time.time()
print(f"Total Time: {end_time - start_time:.4f} seconds")

```

#### MUSE + Logistic Regression

Some weights of the model checkpoint at sentence-transformers/use-cmlm-multilingual were not used when initializing BertModel: ['cls.predictions.bias', 'cls.predictions.transform.LayerNorm.bias', 'cls.predictions.transform.LayerNorm.weight', 'cls.predictions.transform.dense.bias', 'cls.predictions.transform.dense.weight', 'cls.seq\_relationship.bias', 'cls.seq\_relationship.weight']

- This IS expected if you are initializing BertModel from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).
- This IS NOT expected if you are initializing BertModel from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

Encoding Time: 47.9169 seconds

Training Time: 25.3683 seconds

Train Accuracy: 0.7367

Validation Accuracy: 0.7066

Test Accuracy: 0.7016

Total Time: 81.2206 seconds