

hw6-3-wanchanberta

February 14, 2025

1 HOMEWORK 6: TEXT CLASSIFICATION

In this homework, you will create models to classify texts from TRUE call-center. There are two classification tasks: 1. Action Classification: Identify which action the customer would like to take (e.g. enquire, report, cancel) 2. Object Classification: Identify which object the customer is referring to (e.g. payment, true money, internet, roaming)

We will focus only on the Object Classification task for this homework.

In this homework, you are asked to compare different text classification models in terms of accuracy and inference time.

You will need to build 3 different models.

1. A model based on tf-idf
2. A model based on MUSE
3. A model based on wangchanBERTa

You will be asked to submit 3 different files (.pdf from .ipynb) that do the 3 different models. Finally, answer the accuracy and runtime numbers in MCV.

This homework is quite free form, and your answer may vary. We hope that the processing during the course of this assignment will make you think more about the design choices in text classification.

```
[1]: !wget --no-check-certificate https://www.dropbox.com/s/37u83g55p19kvrl/
      ↪ clean-phone-data-for-students.csv
```

```
--2025-02-14 07:03:53-- https://www.dropbox.com/s/37u83g55p19kvrl/clean-phone-
data-for-students.csv
Resolving www.dropbox.com (www.dropbox.com)... 162.125.3.18,
2620:100:6018:18::a27d:312
Connecting to www.dropbox.com (www.dropbox.com)|162.125.3.18|:443... connected.
HTTP request sent, awaiting response... 302 Found
Location: https://www.dropbox.com/sc/fi/8h8hvs9uj6o0524lfe4i/clean-phone-data-
for-students.csv?rlkey=lwv5xbf16jerehmv3lfgq5ue6 [following]
--2025-02-14 07:03:53--
https://www.dropbox.com/sc/fi/8h8hvs9uj6o0524lfe4i/clean-phone-data-for-
students.csv?rlkey=lwv5xbf16jerehmv3lfgq5ue6
Reusing existing connection to www.dropbox.com:443.
HTTP request sent, awaiting response... 302 Found
Location: https://uc279ae1207599e9d9326875105f.dl.dropboxusercontent.com/cd/0/in
```

```

line/CkH0W4KoKEeh3GbbXsnCY1KAfk4Byed22ey5qjE_wN9TVvHj71kfAdPkDI7F7HLgZNUi_w5EgTm
0KEajEE9E01F2yz52dY8VZcxwtbc5o0BsELsNWccsg8xq0JAnJ0bI7SI/file# [following]
--2025-02-14 07:03:54-- https://uc279ae1207599e9d9326875105f.dl.dropboxusercontent.com/cd/0/inline/CkH0W4KoKEeh3GbbXsnCY1KAfk4Byed22ey5qjE_wN9TVvHj71kfAdPkDI7F7HLgZNUi_w5EgTm0KEajEE9E01F2yz52dY8VZcxwtbc5o0BsELsNWccsg8xq0JAnJ0bI7SI/file
Resolving uc279ae1207599e9d9326875105f.dl.dropboxusercontent.com
(uc279ae1207599e9d9326875105f.dl.dropboxusercontent.com)... 162.125.3.15,
2620:100:6018:15::a27d:30f
Connecting to uc279ae1207599e9d9326875105f.dl.dropboxusercontent.com
(uc279ae1207599e9d9326875105f.dl.dropboxusercontent.com)|162.125.3.15|:443...
connected.
HTTP request sent, awaiting response... 200 OK
Length: 2518977 (2.4M) [text/plain]
Saving to: 'clean-phone-data-for-students.csv'

clean-phone-data-fo 100%[=====>] 2.40M --.-KB/s in 0.07s

2025-02-14 07:03:54 (32.7 MB/s) - 'clean-phone-data-for-students.csv' saved
[2518977/2518977]

```

```
[2]: %pip install pythainlp
```

```

Collecting pythainlp
  Downloading pythainlp-5.0.5-py3-none-any.whl.metadata (7.5 kB)
Requirement already satisfied: requests>=2.22.0 in
/usr/local/lib/python3.10/dist-packages (from pythainlp) (2.32.3)
Requirement already satisfied: charset-normalizer<4,>=2 in
/usr/local/lib/python3.10/dist-packages (from requests>=2.22.0->pythainlp)
(3.4.1)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-
packages (from requests>=2.22.0->pythainlp) (3.10)
Requirement already satisfied: urllib3<3,>=1.21.1 in
/usr/local/lib/python3.10/dist-packages (from requests>=2.22.0->pythainlp)
(2.3.0)
Requirement already satisfied: certifi>=2017.4.17 in
/usr/local/lib/python3.10/dist-packages (from requests>=2.22.0->pythainlp)
(2025.1.31)
Downloading pythainlp-5.0.5-py3-none-any.whl (17.9 MB)
17.9/17.9 MB
88.7 MB/s eta 0:00:00:00:0100:01
Installing collected packages: pythainlp
Successfully installed pythainlp-5.0.5
Note: you may need to restart the kernel to use updated packages.

```

1.1 Import Libs

```
[3]: %matplotlib inline
import pandas
import sklearn
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd

from torch.utils.data import Dataset
from IPython.display import display
from collections import defaultdict
from sklearn.metrics import accuracy_score

import warnings
warnings.filterwarnings('ignore')
import time
```

```
[4]: SEED = 42
```

1.2 Loading data

First, we load the data from disk into a Dataframe.

A Dataframe is essentially a table, or 2D-array/Matrix with a name for each column.

```
[5]: data_df = pd.read_csv('clean-phone-data-for-students.csv')
```

Let's preview the data.

```
[6]: # Show the top 5 rows
display(data_df.head())
# Summarize the data
data_df.describe()
```

	Sentence	Utterance	Action	Object
0	<PHONE_NUMBER_REMOVED>	Counte...	enquire	payment
1	internet		enquire	package
2	...	report	suspend	
3	internet	...	enquire	internet
4	...	report	phone_issues	

```
[6]:
```

	Sentence	Utterance	Action	Object
count	16175	16175	16175	
unique	13389	10	33	
top	enquire	service		
freq	97	10377	2525	

1.3 Data cleaning

We call the `DataFrame.describe()` again. Notice that there are 33 unique labels/classes for object and 10 unique labels for action that the model will try to predict. But there are unwanted duplications e.g. `Idd,idd,loyalty_card,Loyalty_card`

Also note that, there are 13389 unique sentence utterances from 16175 utterances. You have to clean that too!

1.4 #TODO 0.1:

- You will have to remove unwanted label duplications as well as duplications in text inputs.
- Also, you will have to trim out unwanted whitespaces from the text inputs.

This shouldn't be too hard, as you have already seen it in the demo.

```
[7]: display(data_df.describe())
display(data_df.Object.unique())
display(data_df.Action.unique())
```

	Sentence Utterance	Action	Object
count	16175	16175	16175
unique	13389	10	33
top	enquire	service	
freq	97	10377	2525

```
array(['payment', 'package', 'suspend', 'internet', 'phone_issues',
       'service', 'nonTrueMove', 'balance', 'detail', 'bill', 'credit',
       'promotion', 'mobile_setting', 'iservice', 'roaming', 'truemoney',
       'information', 'lost_stolen', 'balance_minutes', 'idd',
       'TrueMoney', 'garbage', 'Payment', 'IDD', 'ringtone', 'Idd',
       'rate', 'loyalty_card', 'contact', 'officer', 'Balance', 'Service',
       'Loyalty_card'], dtype=object)

array(['enquire', 'report', 'cancel', 'Enquire', 'buy', 'activate',
       'request', 'Report', 'garbage', 'change'], dtype=object)
```

```
[8]: # TODO 1: Data Cleaning

# Filter cols
cols = ["Sentence Utterance", "Object"]
data_df = data_df[cols]
data_df.columns = ['input', 'raw_label']

# Lowercase: label
data_df['clean_label'] = data_df['raw_label'].str.lower().copy()
data_df.drop('raw_label', axis=1, inplace=True)

# Trim white spaces: input
data_df['input'] = data_df['input'].str.strip()
```

```
# Remove duplicate: input
data_df = data_df.drop_duplicates(subset=['input'], keep='first')

# Display summary
display(data_df.describe())
display(data_df['clean_label'].unique())
```

	input	clean_label
count	13367	13367
unique	13367	26
top	service	
freq	1	2108

```
array(['payment', 'package', 'suspend', 'internet', 'phone_issues',
      'service', 'nontruemove', 'balance', 'detail', 'bill', 'credit',
      'promotion', 'mobile_setting', 'iservice', 'roaming', 'truemoney',
      'information', 'lost_stolen', 'balance_minutes', 'idd', 'garbage',
      'ringtone', 'rate', 'loyalty_card', 'contact', 'officer'],
      dtype=object)
```

Split data into train, validation, and test sets (normally the ratio will be 80:10:10 , respectively). We recommend to use `train_test_split` from `scikit-learn` to split the data into train, validation, test set.

In addition, it should split the data that distribution of the labels in train, validation, test set are similar. There is **stratify** option to handle this issue.

https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.train_test_split.html

Make sure the same data splitting is used for all models.

```
[9]: # Mapping
data = data_df.to_numpy()

unique_label = data_df.clean_label.unique()

label_2_num_map = dict(zip(unique_label, range(len(unique_label))))
num_2_label_map = dict(zip(range(len(unique_label)), unique_label))

print("Create Mappings")
display(num_2_label_map)
display(label_2_num_map)

print("Before Mappings")
display(data[:, 1])
data[:, 1] = np.vectorize(label_2_num_map.get)(data[:, 1]) # Mapping...
print("After Mappings")
display(data[:, 1])

# Trim
```

```

def strip_str(string):
    return string.strip()
print("Before")
print(data)
data[:,0] = np.vectorize(strip_str)(data[:,0]) # Trimming...
print("After")
print(data)

```

Create Mappings

```

{0: 'payment',
 1: 'package',
 2: 'suspend',
 3: 'internet',
 4: 'phone_issues',
 5: 'service',
 6: 'nontruemove',
 7: 'balance',
 8: 'detail',
 9: 'bill',
10: 'credit',
11: 'promotion',
12: 'mobile_setting',
13: 'iservice',
14: 'roaming',
15: 'truemoney',
16: 'information',
17: 'lost_stolen',
18: 'balance_minutes',
19: 'idd',
20: 'garbage',
21: 'ringtone',
22: 'rate',
23: 'loyalty_card',
24: 'contact',
25: 'officer'}

{'payment': 0,
 'package': 1,
 'suspend': 2,
 'internet': 3,
 'phone_issues': 4,
 'service': 5,
 'nontruemove': 6,
 'balance': 7,
 'detail': 8,
 'bill': 9,
 'credit': 10,
 'promotion': 11,

```

```

'mobile_setting': 12,
'iservice': 13,
'roaming': 14,
'truemoney': 15,
'information': 16,
'lost_stolen': 17,
'balance_minutes': 18,
'idd': 19,
'garbage': 20,
'ringtone': 21,
'rate': 22,
'loyalty_card': 23,
'contact': 24,
'officer': 25}

```

Before Mappings

```

array(['payment', 'package', 'suspend', ..., 'balance', 'balance',
      'package'], dtype=object)

```

After Mappings

```

array([0, 1, 2, ..., 7, 7, 1], dtype=object)

```

Before

```

[['<PHONE_NUMBER_REMOVED>      Counter Services      3276.25
  3057.79      '
  0]
['internet      ' 1]
['      ' 2]
...
['      ' 7]
['      ' 7]
['      ' 1]]

```

After

```

[['<PHONE_NUMBER_REMOVED>      Counter Services      3276.25
  3057.79      '
  0]
['internet      ' 1]
['      ' 2]
...
['      ' 7]
['      ' 7]
['      ' 1]]

```

```

[10]: def filter_data(df):
      class_counts = df["label"].value_counts()
      valid_classes = class_counts[class_counts >= 10].index
      invalid_classes = class_counts[class_counts < 10].index

```

```

correct_classes = [i for i, _ in enumerate(valid_classes.sort_values())]
change_map = dict(zip(valid_classes.sort_values(), correct_classes))

# drop invalid classes
df = df[df["label"].isin(valid_classes)]
for num in invalid_classes:
    label = num_2_label_map[num]
    del num_2_label_map[num]
    del label_2_num_map[label]

# update change
df["label"] = df["label"].apply(lambda x: change_map.get(x, -1))
for old_num, new_num in change_map.items():
    num_2_label_map[new_num] = num_2_label_map.pop(old_num)
    label_2_num_map[num_2_label_map[new_num]] = new_num
return df

df = pd.DataFrame(data, columns=['input', 'label'])
df_filtered = filter_data(df)

```

```

[11]: # TODO: Split data
from sklearn.model_selection import train_test_split

def split_data(data_df, random_state=SEED):
    """split_data splits the data into train:validation:test=8:1:1 sets."""
    X = data_df["input"]
    y = data_df["label"]

    # First split: Train (80%) and Temp (20%)
    X_train, X_temp, y_train, y_temp = train_test_split(
        X, y, test_size=0.20, stratify=y, random_state=random_state
    )

    # Second split: Validation (10%) and Test (10%)
    X_val, X_test, y_val, y_test = train_test_split(
        X_temp, y_temp, test_size=0.50, stratify=y_temp,
        random_state=random_state
    )

    # Display dataset sizes
    print(f"Train size: {len(X_train)}")
    print(f"Validation size: {len(X_val)}")
    print(f"Test size: {len(X_test)}")

    return X_train, X_val, X_test, y_train, y_val, y_test

# Split

```



```
X_train, X_val, X_test, y_train, y_val, y_test = split_data(df_filtered)
```

Train size: 10690

Validation size: 1336

Test size: 1337

2 Model 3 WangchanBERTa

We ask you to train a WangchanBERTa-based model.

We recommend you use the thaixtransformers fork (which we used in the PoS homework).
<https://github.com/PyThaiNLP/thaixtransformers>

The structure of the code will be very similar to the PoS homework. You will also find the huggingface [tutorial](#) useful. Or you can also add a softmax layer by yourself just like in the previous homework.

Which WangchanBERTa model will you use? Why? (Don't forget to clean your text accordingly).

Ans: airesearch/wangchanberta-base-att-spm-uncased because - specifically trained for Thai text. - SentencePiece tokenization, which is ideal for Thai, unlike space-based tokenization. - It has state-of-the-art performance for Thai text classification tasks.

```
[12]: # Data preprocessing
      for i in range(len(data)):
          data[i][0] = data[i][0].replace(' ', " ")
```

```
[13]: from kaggle_secrets import UserSecretsClient
      user_secrets = UserSecretsClient()
      secret_value_0 = user_secrets.get_secret("WANDB_API_KEY")

      import wandb
      wandb.login(key=secret_value_0)
```

wandb: Using wandb-core as the SDK backend. Please refer to <https://wandb.me/wandb-core> for more information.

wandb: Currently logged in as: **pupipat-sk**

(**pupipatsk**). Use `wandb login --relogin` to force relogin

wandb: **WARNING** If you're specifying your api key in code, ensure this code is not shared publicly.

wandb: **WARNING** Consider setting the WANDB_API_KEY environment variable, or running `wandb login` from the command line.

wandb: Appending key for api.wandb.ai to your netrc file:
/root/.netrc

```
[13]: True
```

```
[14]: import torch
      def find_device() -> str:
          device = "cpu"
```

```

    if torch.cuda.is_available():
        device = "cuda"
    elif torch.backends.mps.is_available():
        device = "mps"
    else:
        device = "cpu"
    print(f"device: {device}")
    return device
device = find_device()

```

device: cuda

```

[ ]: import numpy as np
import torch
from transformers import AutoTokenizer, AutoModelForSequenceClassification,
    ↪Trainer, TrainingArguments, DataCollatorWithPadding
from datasets import Dataset
from sklearn.metrics import accuracy_score

start_time = time.time()

# Define the WangchanBERTa model
MODEL_NAME = "airesearch/wangchanberta-base-att-spm-uncased"
tokenizer = AutoTokenizer.from_pretrained(MODEL_NAME)

# Tokenization function with dynamic padding
def tokenize_function(examples):
    return tokenizer(examples["text"], truncation=True) # Removed "max_length"
    ↪and fixed padding

# Convert data into Hugging Face Dataset format
train_data = Dataset.from_dict({"text": X_train.tolist(), "label": y_train.
    ↪tolist()})
val_data = Dataset.from_dict({"text": X_val.tolist(), "label": y_val.tolist()})
test_data = Dataset.from_dict({"text": X_test.tolist(), "label": y_test.
    ↪tolist()})

# Tokenize datasets
train_dataset = train_data.map(tokenize_function, batched=True)
val_dataset = val_data.map(tokenize_function, batched=True)
test_dataset = test_data.map(tokenize_function, batched=True)

# Load pre-trained WangchanBERTa model for classification
model = AutoModelForSequenceClassification.from_pretrained(MODEL_NAME,
    ↪
    ↪num_labels=len(num_2_label_map),

```

```

↪id2label=num_2_label_map,

↪label2id=label_2_num_map)
model.to(device)

# Define training arguments
EPOCHS = 5
BATCH_SIZE = 32
training_args = TrainingArguments(
    output_dir="./results",
    learning_rate=2e-5,
    per_device_train_batch_size=BATCH_SIZE,
    per_device_eval_batch_size=BATCH_SIZE,
    num_train_epochs=EPOCHS,
    weight_decay=0.01,
    evaluation_strategy="epoch",
    save_strategy="epoch",
    load_best_model_at_end=True,
    push_to_hub=False,
    logging_dir="./logs",
    logging_steps=500,
)

# Compute accuracy
def compute_metrics(eval_pred):
    predictions, labels = eval_pred
    preds = np.argmax(predictions, axis=1)
    return {"accuracy": accuracy_score(labels, preds)}

# Use dynamic padding
data_collator = DataCollatorWithPadding(tokenizer=tokenizer)

# Define trainer with dynamic padding
trainer = Trainer(
    model=model,
    args=training_args,
    train_dataset=train_dataset,
    eval_dataset=val_dataset,
    compute_metrics=compute_metrics,
    data_collator=data_collator # Enables dynamic padding
)

# Train model...
print("Start training...")
trainer.train()
print("Finish training.")

```

```

# Evaluate model on validation and test sets
train_results = trainer.evaluate(train_dataset)
val_results = trainer.evaluate(val_dataset)
test_results = trainer.evaluate(test_dataset)

print(f"Train Accuracy: {train_results['eval_accuracy']:.4f}")
print(f"Validation Accuracy: {val_results['eval_accuracy']:.4f}")
print(f"Test Accuracy: {test_results['eval_accuracy']:.4f}")
end_time = time.time()
print(f"Total Time: {end_time - start_time:.4f} seconds")

```

tokenizer_config.json: 0%| | 0.00/282 [00:00<?, ?B/s]

config.json: 0%| | 0.00/546 [00:00<?, ?B/s]

sentencepiece.bpe.model: 0%| | 0.00/905k [00:00<?, ?B/s]

Map: 0%| | 0/10690 [00:00<?, ? examples/s]

Asking to truncate to max_length but no maximum length is provided and the model has no predefined maximum length. Default to no truncation.

Map: 0%| | 0/1336 [00:00<?, ? examples/s]

Map: 0%| | 0/1337 [00:00<?, ? examples/s]

model.safetensors: 0%| | 0.00/423M [00:00<?, ?B/s]

Some weights of CamembertForSequenceClassification were not initialized from the model checkpoint at airesearch/wangchanberta-base-att-spm-uncased and are newly initialized: ['classifier.dense.bias', 'classifier.dense.weight', 'classifier.out_proj.bias', 'classifier.out_proj.weight']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

wandb: WARNING The `run_name` is currently set to the same value as `TrainingArguments.output_dir`. If this was not intended, please specify a different run name by setting the `TrainingArguments.run_name` parameter.

Start training...

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

```
[17]: # Evaluate model on validation and test sets
train_results = trainer.evaluate(train_dataset)
val_results = trainer.evaluate(val_dataset)
test_results = trainer.evaluate(test_dataset)

print(f"Train Accuracy: {train_results['eval_accuracy']:.4f}")
print(f"Validation Accuracy: {val_results['eval_accuracy']:.4f}")
print(f"Test Accuracy: {test_results['eval_accuracy']:.4f}")
```

Train Accuracy: 0.7889
Validation Accuracy: 0.7657
Test Accuracy: 0.7450

3 Comparison

After you have completed the 3 models, compare the accuracy, ease of implementation, and inference speed (from cleaning, tokenization, till model compute) between the three models in mycourseville.

Model 1: TfidfVectorizer + Logistic Regression + pythainlp.word_tokenize - Training time: 1.5351 seconds - Train Accuracy: 0.7675 - Validation Accuracy: 0.6894 - Test Accuracy: 0.6911

Model 2: MUSE + Logistic Regression - Encoding Time: 47.9169 seconds - Training Time: 25.3683 seconds - Train Accuracy: 0.7367 - Validation Accuracy: 0.7066 - Test Accuracy: 0.7016 - Total Time: 81.2206 seconds

Model 3: WangchanBERTa - Train Accuracy: 0.7889 - Validation Accuracy: 0.7657 - Test Accuracy: 0.7450 - Total Time: 5 min 10 sec

Based on the performance of the three models, which one do you think is best for this use case (Callcenter Chatbot)?

Answer: WangchanBERTa

because we want CallcenterChatbot that doesn't need that fast respond(compare to current answer bot when we call for customer service) - Highest accuracy, callcenter question can be repetitive question - Slow inference time, but can speedup by increase computation - Better generalization - Deep contextual understanding for handles nuances and variations in Thai language